

ML & Neural Network Tutorial

ANALS 2023

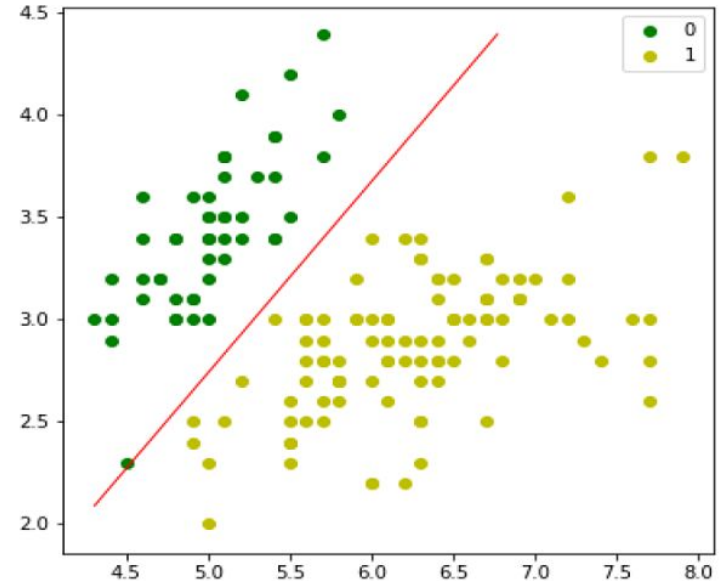
Nripesh Parajuli, PhD

Different types of ML algorithms

- KNNs
- Logistic regression
- Support vector machines (SVC / SVR)
- Random Forest
- Boosted trees
- Neural networks / Deep Learning

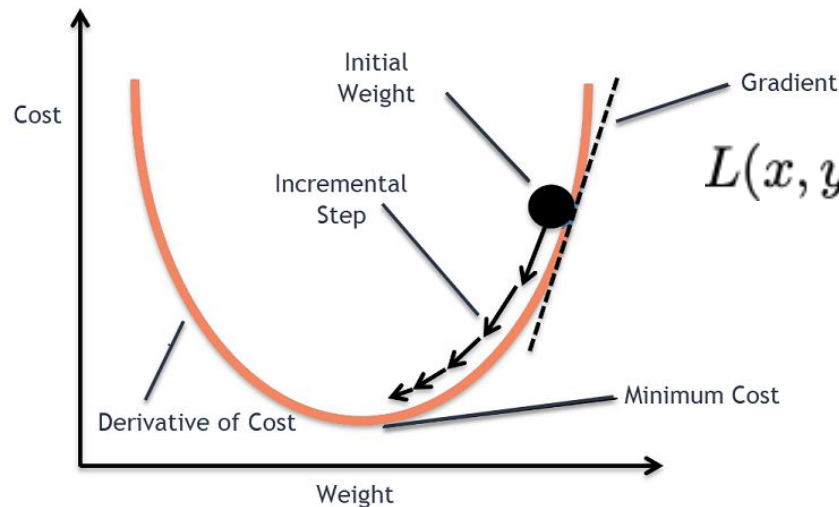
Linear model: Logistic regression

$$p(x) = \frac{1}{1 + e^{-\beta x}}$$



Finding the optimal parameters

- How to find the optimal beta?
- For 1-2 dimensions, we can do trial and error, but not scalable.
- We minimize the following loss function using gradient descent:



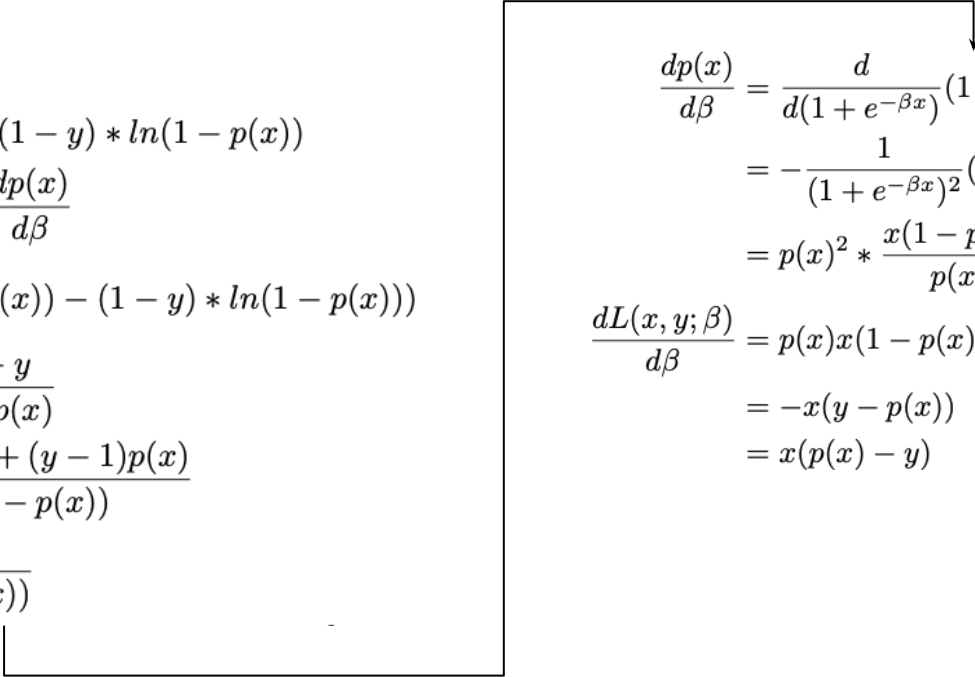
$$L(x, y) = -y * \log(p(x)) - (1 - y) * \log(1 - p(x))$$

$$\nabla_{\beta} L = \frac{dL(x, y; \beta)}{d\beta}$$

$$\beta' = \beta - LR * \nabla_{\beta} L$$

Gradient computation*

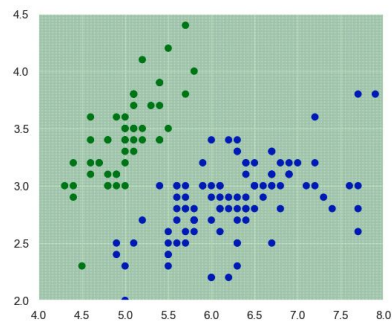
$$\begin{aligned}p(x) &= \frac{1}{1 + e^{-\beta x}} \\L(x, y; \beta) &= -y \ln(p(x)) - (1 - y) * \ln(1 - p(x)) \\\frac{dL(x, y; \beta)}{d\beta} &= \frac{dL(x, y; \beta)}{dp(x)} * \frac{dp(x)}{d\beta} \\\frac{dL(x, y; \beta)}{dp(x)} &= -\frac{d}{dp(x)}(y \ln(p(x)) - (1 - y) * \ln(1 - p(x))) \\&= -\frac{y}{p(x)} + \frac{1 - y}{1 - p(x)} \\&= -\frac{y(1 - p(x)) + (y - 1)p(x)}{p(x)(1 - p(x))} \\&= -\frac{y - p(x)}{p(x)(1 - p(x))}\end{aligned}$$


$$\begin{aligned}\frac{dp(x)}{d\beta} &= \frac{d}{d(1 + e^{-\beta x})}(1 + e^{-\beta x})^{-1} * \frac{d(1 + e^{-\beta x})}{d\beta} \\&= -\frac{1}{(1 + e^{-\beta x})^2}(-xe^{-\beta x}) \\&= p(x)^2 * \frac{x(1 - p(x))}{p(x)} = p(x)x(1 - p(x)) \\\frac{dL(x, y; \beta)}{d\beta} &= p(x)x(1 - p(x)) * -\frac{(y - p(x))}{p(x)(1 - p(x))} \\&= -x(y - p(x)) \\&= x(p(x) - y)\end{aligned}$$

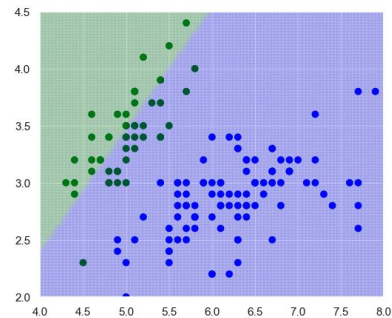
Gradient descent

- Initialize **beta** with some random values or **zero**.
- For each step:
 - Compute the gradient of the loss function at that value of beta, call it **grad**.
 - Update **beta** = **beta** - **LR** * **grad**. Where **LR** is the learning.
 - Terminate if the magnitude of the change (grad) is too small or after **T** number of iterations.
- Learning rate (**LR**) can be adapted as steps progress. Higher in the beginning and slower later.
- **T** also needs to be adapted depending on the problem.
- Gradient can be computed directly (exactly) or analytically using a small perturbation around the variable.

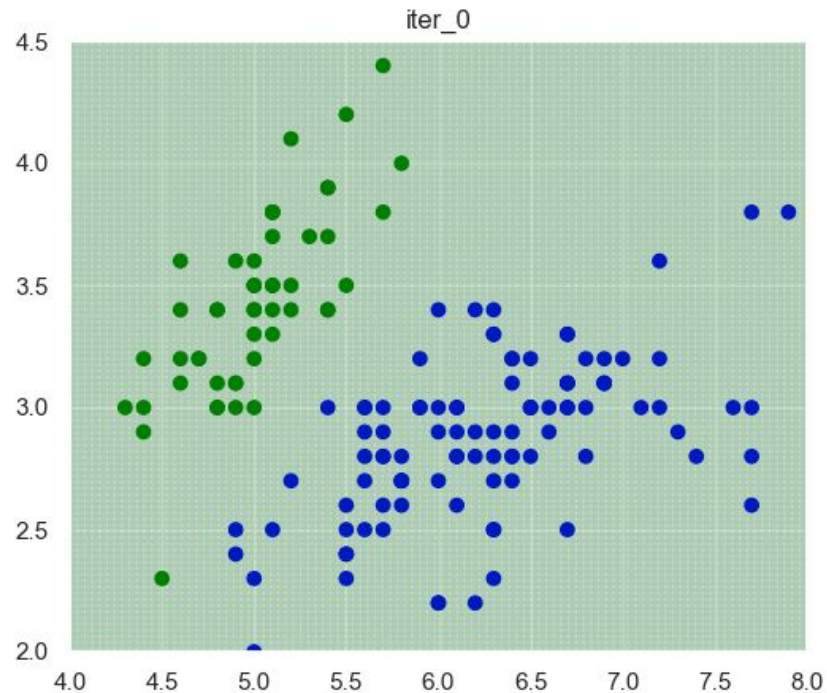
Gradient descent



iter=0

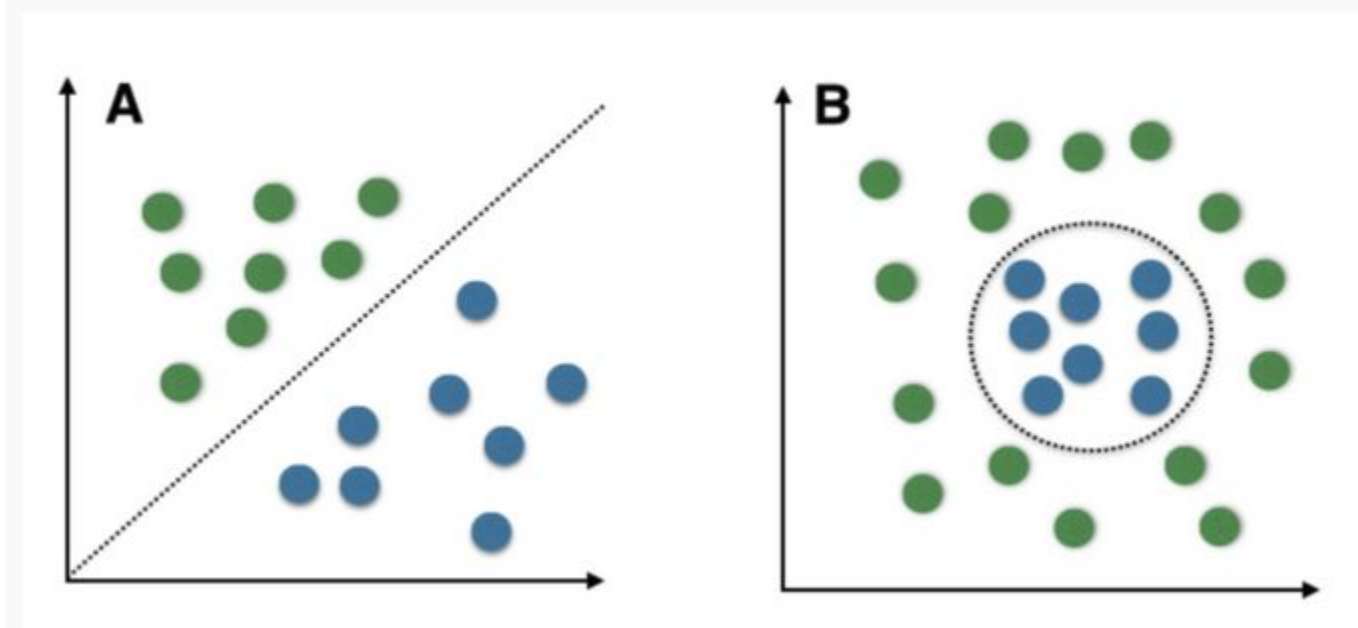


iter=150

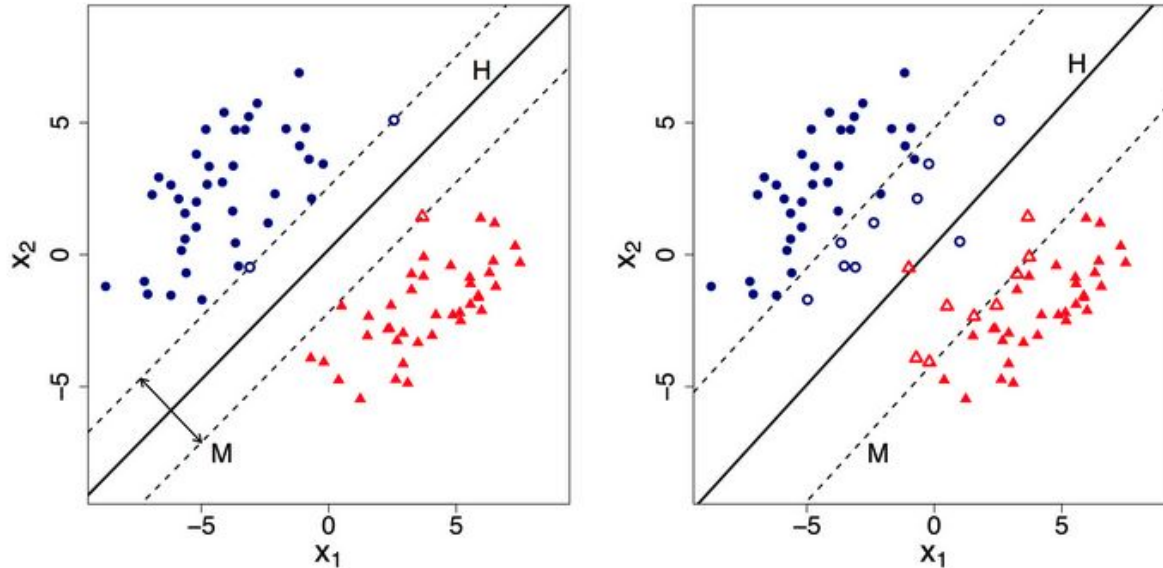


Exercise: implementing logistic regression w gradient descent

Importance of non-linearity



Support vector machines



Hard margin (left) and soft margin (right) SVMs.

Tree based methods

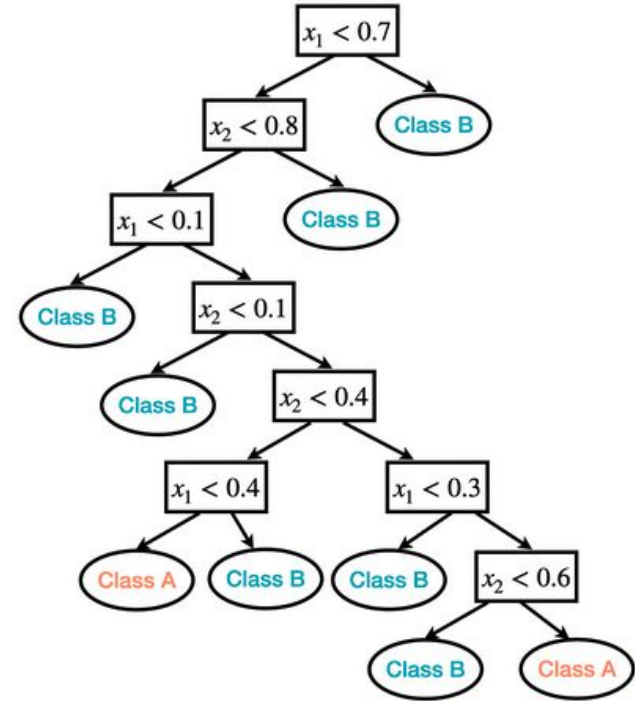
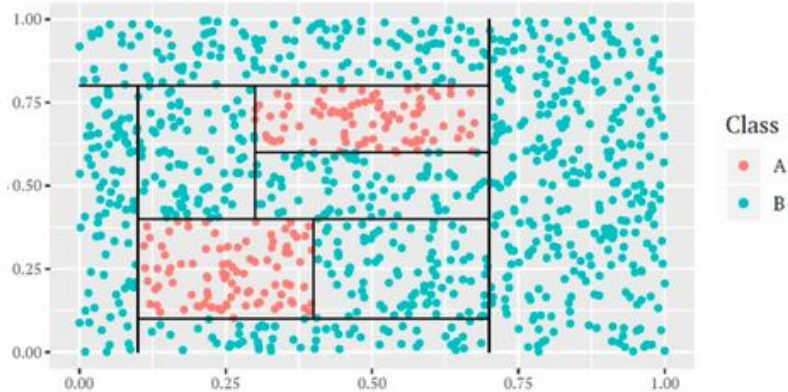


Figure 13.1: An example of classification tree based one two predictors.

Neural network: backpropagation

- Gradient descent is used to optimize the parameters for neural networks as well.
- More specifically, **stochastic gradient descent**, is used:
 - This just means gradient descent is done in a 'random' portion of the data called the **batch**.
- Backpropagation is used to propagate the gradient back to the layers and update weights.

Summary: the equations of backpropagation

$$\delta^L = \nabla_a C \odot \sigma'(z^L) \quad (\text{BP1})$$

$$\delta^l = ((w^{l+1})^T \delta^{l+1}) \odot \sigma'(z^l) \quad (\text{BP2})$$

$$\frac{\partial C}{\partial b_j^l} = \delta_j^l \quad (\text{BP3})$$

$$\frac{\partial C}{\partial w_{jk}^l} = a_k^{l-1} \delta_j^l \quad (\text{BP4})$$

Neural network: backpropagation

Feedforward part:

for i in range(len(layers)):

$$a(i) = o(i-1).w[i]$$

$$o(i) = \frac{1}{1 + e^{-a(i)}}$$

$$= \text{sigmoid}(a(i))$$

Backpropagate part:

for the output layer :

$$\text{error}[L] = o[L] - y$$

$$\text{grad}[L] = \text{error} * \text{sigmoid_deriv}(o[L])$$

for other layers :

$$\text{error}[i] = \text{grad}[i+1].W[i]$$

$$\text{grad}[i] = \text{error}[i] * \text{sigmoid_deriv}(o[i])$$

Exercise: looking at different components of the backpropagation algorithm.